

## Outline

- Optimization
- CNN basics
- Examples of CNN Architectures
- Applications of CNN

3

## Optimization


- Stochastic Gradient Descent
- Momentum Method
- Adaptive Learning Methods
- (AdaGrad, RMSProp, Adam)

Broadly applicable for many ML methods

- Batch Normalization ← specific to neural nets

4

## Stochastic Gradient Descent



Slide credit: Shubhendu Trivedi, Risi Kondor 5


## Limitations of GD and SGD

- GD and SGD suffer in the following scenarios:
  - Error surface has high curvature
  - We get small but consistent gradients
  - The stochastic gradients are very noisy (for SGD)

Slide credit: Shubhendu Trivedi, Risi Kondor 6

## Momentum

- The Momentum method is a method to accelerate learning using GD or SGD




- Illustration of how momentum traverses such an error surface better compared to gradient descent
- [Visualization of GD with momentum](#)

Slide credit: Shubhendu Trivedi, Risi Kondor 7

## Momentum

$$\mathbf{v}_{\text{new}} = \alpha \mathbf{v}_{\text{prev}} - \epsilon \nabla_{\theta} \left( \mathcal{L} \left( f(\mathbf{x}^{(i)}; \theta), \mathbf{y}^{(i)} \right) \right)$$



[Visualization of GD with momentum](#)

Slide credit: Shubhendu Trivedi, Risi Kondor 8

## Adaptive Learning Methods: Motivation

- Until now we assigned the same learning rate to all features
  - If the features vary in importance and frequency, is this a good idea?
  - Probably not!
- Popular methods for adaptive learning rates:
  - AdaGrad, RMSProp, Adam, etc.
- High-level idea:
  - Discount the learning rate for each parameter by dividing with the “amplitude” (running average) of the gradient for that parameter.
  - Training is more stable and robust to noisy gradients and the choice of initial learning rates (e.g., even a reasonably large initial learning rate works)

9

## Optimization with SGD and its variants

$$\text{SGD: } \theta \leftarrow \theta - \epsilon \hat{\mathbf{g}}$$

$$\text{Momentum: } \mathbf{v} \leftarrow \alpha \mathbf{v} - \epsilon \hat{\mathbf{g}} \text{ then } \theta \leftarrow \theta + \mathbf{v}$$

$$\text{AdaGrad: } \mathbf{r} \leftarrow \mathbf{r} + \hat{\mathbf{g}} \odot \hat{\mathbf{g}} \text{ then } \Delta \theta \leftarrow \frac{\epsilon}{\delta + \sqrt{\mathbf{r}}} \odot \hat{\mathbf{g}} \text{ then } \theta \leftarrow \theta + \Delta \theta$$

$$\text{RMSProp: } \mathbf{r} \leftarrow \rho \mathbf{r} + (1 - \rho) \hat{\mathbf{g}} \odot \hat{\mathbf{g}} \text{ then } \Delta \theta \leftarrow \frac{\epsilon}{\delta + \sqrt{\mathbf{r}}} \odot \hat{\mathbf{g}} \text{ then } \theta \leftarrow \theta + \Delta \theta$$

$$\text{ADAM: } \mathbf{s} \leftarrow \rho_1 \mathbf{s} + (1 - \rho_1) \hat{\mathbf{g}} \\ \mathbf{r} \leftarrow \rho_2 \mathbf{r} + (1 - \rho_2) \hat{\mathbf{g}} \odot \hat{\mathbf{g}}$$

$$\hat{\mathbf{s}} \leftarrow \frac{\mathbf{s}}{1 - \rho_1^t}, \hat{\mathbf{r}} \leftarrow \frac{\mathbf{r}}{1 - \rho_2^t} \text{ then } \Delta \theta = -\epsilon \frac{\hat{\mathbf{s}}}{\sqrt{\hat{\mathbf{r}} + \delta}} \text{ then } \theta \leftarrow \theta + \Delta \theta$$

- Visualization:
  - [Visualization of optimizers](#)
  - [Distill: why momentum really works](#)

Slide credit: Shubhendu Trivedi, Risi Kondor 16

# Batch Normalization



- **BN is specific to neural networks**
- We have a recipe to compute gradients (backpropagation) and update all parameters (SGD, adaptive learning rate methods, etc.)
- **Challenge: Internal Covariate Shift**
  - Implicit assumption during backpropagation: layer inputs remain unchanged.
  - Reality: Simultaneous updates across layers alter distributions.
  - Consequences:
    - Shifts in activation distributions across layers.
    - Training instability: unstable training, longer convergence times, suboptimal performance.
- **Insights: Addressing internal covariate shift can significantly enhance training efficiency.**

17

## Batch normalization standardizes inputs to each layer...

- ...which helps to stabilize training
- Consider standardizing the input to the input layer, i.e. our data
  - For minibatch mean  $\mu_B$  and standard deviation  $\sigma_B^2$ , we can normalize each input  $x_i$  to
 
$$\hat{x}_i \leftarrow \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} \quad (\epsilon \text{ added for numerical stability})$$
 which shifts our input to a distribution w/ mean = 0 and std dev = 1
  - Output of batch normalization (with  $\gamma, \beta$  being learnable parameters for the BN layer):
 
$$y_i \leftarrow \gamma \hat{x}_i + \beta$$
- Batch normalization extends this idea to the input of every layer, not just the input layer
  - But after going through previous layers and activations, the input does not necessarily reflect the original input distribution
  - To reflect the true distribution of the data, keep track of scale  $\gamma$  and shift  $\beta$  for each weight

19

## Batch Normalization: forward prop

- Normalize distribution of each input feature in each layer across each minibatch to  $N(0, 1)$
- Learn the scale and shift
- After training, at test time: Use running averages of  $\mu$  and  $\sigma$  collected during training, use these for evaluating new input  $x$

**Input:** Values of  $x$  over a mini-batch:  $B = \{x_1 \dots x_m\}$ ;  
Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_B \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_B^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

Ioffe, S., & Szegedy, C. (2015). Batch normalization: Accelerating deep network training by reducing internal covariate shift.

20

## Batch Normalization: backpropagation

- Differentiable via chain rule

$$\frac{\partial \ell}{\partial \hat{x}_i} = \frac{\partial \ell}{\partial y_i} \cdot \gamma$$

$$\frac{\partial \ell}{\partial \sigma_B^2} = \sum_{i=1}^m \frac{\partial \ell}{\partial \hat{x}_i} \cdot (x_i - \mu_B) \cdot \frac{-1}{2} (\sigma_B^2 + \epsilon)^{-3/2}$$

$$\frac{\partial \ell}{\partial \sigma_B} = \sum_{i=1}^m \frac{\partial \ell}{\partial \hat{x}_i} \cdot \frac{-1}{\sqrt{\sigma_B^2 + \epsilon}}$$

$$\frac{\partial \ell}{\partial x_i} = \frac{\partial \ell}{\partial \hat{x}_i} \cdot \frac{1}{\sqrt{\sigma_B^2 + \epsilon}} + \frac{\partial \ell}{\partial \sigma_B^2} \cdot \frac{2(x_i - \mu_B)}{m} + \frac{\partial \ell}{\partial \sigma_B} \cdot \frac{1}{m}$$

$$\frac{\partial \ell}{\partial \gamma} = \sum_{i=1}^m \frac{\partial \ell}{\partial y_i} \cdot \hat{x}_i$$

$$\frac{\partial \ell}{\partial \beta} = \sum_{i=1}^m \frac{\partial \ell}{\partial y_i}$$

Ioffe, S., & Szegedy, C. (2015). Batch normalization: Accelerating deep network training by reducing internal covariate shift.

21